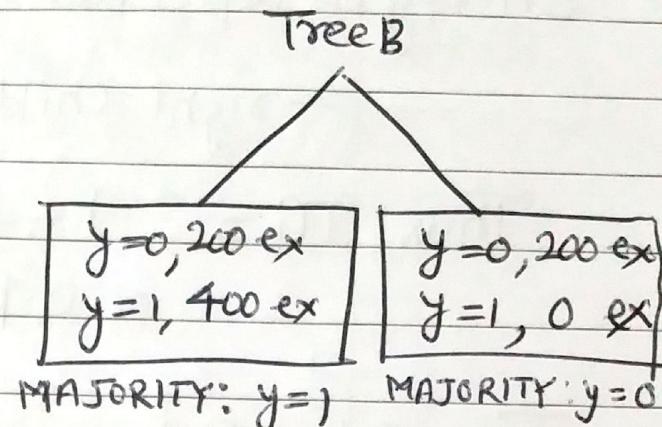
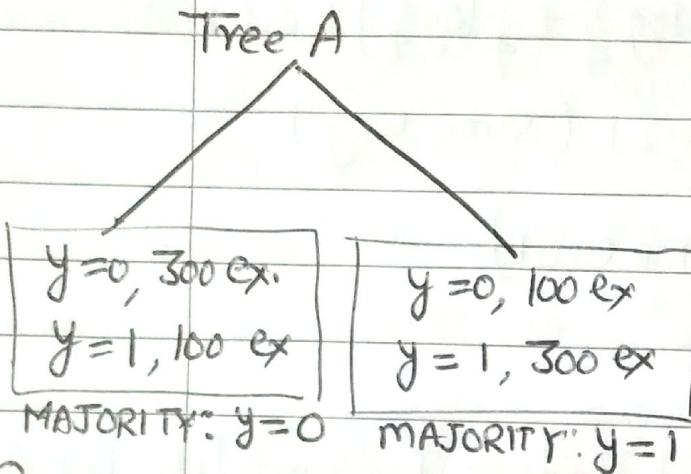


Homework 1 (Problem 1)



- ① Tree A: 200 misclassifications (100 ex. with $y=1$ are misclassified in left branch, and 100 ex. with $y=0$ are misclassified in right branch)

$$\text{Tree A misclassification rate} = \frac{200}{800} = 0.25$$

- Tree B: 200 misclassifications (200 ex. with $y=0$ are misclassified in left branch, and none of the examples are misclassified in the right branch)

$$\text{Tree B misclassification rate} = \frac{200}{800} = 0.25$$

So the misclassification rate is equal for both.

- ② Tree A: Entropy of root node = $-\left[\frac{100}{800} \log \frac{100}{800} + \frac{400}{800} \log \frac{400}{800}\right] = 1$

$$\text{Entropy of left child} = -\left[\frac{3}{4}\log\frac{3}{4} + \frac{1}{4}\log\frac{1}{4}\right] = 0.81$$

right child = 0.81 (same way)

$$\text{Thus } IG = 1 - \frac{1}{2}[0.81 + 0.81] \\ \approx 0.19$$

Tree B: Entropy of root node = 1 (just like Tree A)

$$\text{Entropy of left child} = -\left[\frac{2}{3}\log\frac{2}{3} + \frac{1}{3}\log\frac{1}{3}\right] \approx 0.92$$

right child = 0 (pure node, no ex. with $y=1$)

$$\text{Thus } IG = 1 - \left[\frac{3}{4} \times 0.92 + \frac{1}{4} \times 0\right] \approx 0.31$$

Thus Tree B has a higher IG.

③ Indeed, the answers are different. As per entropy/IG criterion, Tree B is preferred.

It makes sense because IG based criterion is sensitive to node probabilities as opposed to misclassification rate which is only based on a "hard" classification as 0/1.

Problem 2

In the limit of infinite data, and since we have a noise-free setting, each input location has ~~one~~ a training example with its correct label.

Since the training set is infinite, a new test input will fall exactly on one of the training inputs and will get the same label (since we are using a one-nearest neighbor).

Problem 3

$$\begin{aligned} f(x_*) &= x_*^T w \\ &= x_*^T (x^T x)^{-1} x^T y \\ &= x_*^T (x^T x)^{-1} \sum_{n=1}^N x_n y_n \\ &= \sum_{n=1}^N x_*^T \underbrace{\sum_{n=1}^{-1} x_n y_n}_{w_n} \\ &\equiv \sum_{n=1}^N w_n y_n \end{aligned}$$

11

Note the form of each W_n :

$$W_n = \mathbf{x}_*^T \Sigma \mathbf{x}_n$$

This is a general form of an inner product similarity between the test input \mathbf{x}_* and training input \mathbf{x}_n , modulated by the matrix $\Sigma = \mathbf{X}^T \mathbf{X}$, which depends on the entire training data.

In contrast, in ~~KNN~~ the weighted version of KNN, the similarities only depend on the test input and the nearest neighbor (e.g. $W_n = \mathbf{x}_*^T \mathbf{x}_n$ or $W_n = \frac{1}{\|\mathbf{x}_* - \mathbf{x}_n\|}$).

~~Problem~~ Problem 4

To have a different amount of regularization on each component of W , we can modify the regularizer (ℓ_2) from $\lambda W^T W = \sum_{d=1}^D \lambda w_d^2$ to $\sum_{d=1}^D \lambda_d w_d^2$ where λ_d denotes the regularization hyperparameter for component w_d of the weight vector W .

Note that we can also write it ~~as~~ in terms of vector and matrix as follows:

$\sum_{d=1}^D \lambda_d w_d^2 = w^T \Lambda w$ where Λ is
 a diagonal matrix of size $D \times D$ with
 $\{\lambda_d\}_{d=1}^D$ along its diagonals.

The full objective will be:

$$f(w) = \sum_{n=1}^N (y_n - w^T x_n)^2 + w^T \Lambda w$$

Solving this for w is almost identical to the way we have seen in the class, so I'll skip the full derivation. The solution will be

$$\hat{w} = (X^T X + \Lambda)^{-1} X^T y$$

(An aside): Probabilistic Version: Just as $N(0, \lambda^{-1} I)$ prior on w corresponds to an L_2 regularizer, changing the prior to $N(0, \Lambda^{-1})$ will correspond to the above regularizer with component-wise regularization.

Problem 5

$$L(w) = \text{trace} [(Y - Xw)^T (Y - Xw)]$$

with $W = BS$, the objective will be

$$L(B, S) = \text{trace} [(Y - XBS)^T (Y - XBS)]$$

If B is known, the objective is convex in S
(note that it is just like a quadratic function of S)

To get the optimal S , we take the derivative
of $L(S)$ with respect to S (note that since B
is assumed known, I wrote $L(B, S)$ as just $L(S)$)
and set it to zero. Just like the linear
regression case, it is possible to separate S and
get a closed form expression.

$$\frac{\partial}{\partial S} L(S) = 0$$

$$\Rightarrow \frac{\partial}{\partial S} \text{trace} [(Y - \tilde{X}S)^T (Y - \tilde{X}S)] = 0$$

(where I have denoted $\tilde{X} = XB$)

For the above derivative, you can either
directly use the expression of derivative
of trace(quadratic form) from matrix

Cookbook, or expand it as

$$\text{trace}[(Y - \tilde{X}S)^T(Y - \tilde{X}S)]$$

$$= \text{trace}[(Y^T - S^T \tilde{X}^T)(Y - \tilde{X}S)]$$

Multiply the terms and then take derivatives w.r.t. S .

In either case, the final expression will turn out to be (skipping the equations of derivative, since we already did it in PS-1 and also on Piazza).

$$\hat{S} = (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T Y \quad (\text{with } \tilde{X} = XB)$$

which is exactly in the same form as the solution of standard multi-output regression (PS-1) with X replaced by $\tilde{X} = XB$ (it is like reducing the dimensions of X from D to K by multiplying with B)

Bonus Problem (not for credit): In this case, we need to solve for both B and S .

We can use alternating optimization (assume B fixed, solve for S ; assume S fixed, solve for B).

Expression for S , given B will be the same as above. Expression for B , given S will be

$$\hat{B} = (X^T X)^{-1} X^T Y S^T (S S^T)^{-1}$$

Another "Shorthand" to find the solution:

As per the multiooutput regression model

$$XW = Y$$

$$XBS = Y$$

$$X^T X BS = X^T Y$$

$$\underbrace{B^T X^T X}_{\substack{\text{square matrix} \\ \text{now}}} BS = B^T X^T Y \quad (\text{multiplied by } B^T \text{ so that we can invert it in the next step})$$

$$S = (B^T X^T X B)^{-1} B^T X^T Y$$

$$S = (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T Y$$

We can do the same for B (assuming S known)

Starting with

$$X^T X BS = X^T Y$$

$$X^T X B S S^T = X^T Y S^T$$

Premultiply by $(X^T X)^{-1}$ and postmultiply by $(S S^T)^{-1}$

$$B = (X^T X)^{-1} X^T Y S^T (S S^T)^{-1}$$

Note that in this approach, we didn't solve the optimization problem of minimizing the trace term, but rather solved it as linear system of equations (they are equivalent).